

## Exploratory Data Analysis (EDA) - Real State Analysis

*#importing libraries*

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
%matplotlib inline
```

```
df = pd.read_csv('combined_csv.csv')
df.head(10)
```

	SALE TYPE	SOLD DATE	PROPERTY TYPE	\
0	PAST SALE	September-15-2022	Condo/Co-op	
1	PAST SALE	August-31-2022	Condo/Co-op	
2	PAST SALE	September-30-2022	Single Family Residential	
3	PAST SALE	September-28-2022	Vacant Land	
4	PAST SALE	August-25-2022	Condo/Co-op	
5	PAST SALE	October-12-2022	Vacant Land	
6	PAST SALE	August-31-2022	Condo/Co-op	
7	PAST SALE	September-30-2022	Single Family Residential	
8	PAST SALE	August-25-2022	Condo/Co-op	
9	PAST SALE	October-12-2022	Vacant Land	

	ADDRESS	CITY	STATE OR PROVINCE	\
0	84-770 Kili Dr #1440	Waianae	HI	
1	85-175 Farrington Hwy Unit C112	Waianae	HI	
2	16-743 Wao Kele Rd (road G)	Mountain View	HI	
3	Lot Unit 7	Papaikou	HI	
4	4999 Kahala Ave #271	Honolulu	HI	
5		Mountain View	HI	
6	85-175 Farrington Hwy Unit C112	Waianae	HI	
7	16-743 Wao Kele Rd (road G)	Mountain View	HI	
8	4999 Kahala Ave #271	Honolulu	HI	
9	-	Mountain View	HI	

	ZIP OR POSTAL CODE	PRICE	BEDS	BATHS	...	STATUS	\
0	96792	159000	0.0	1.0	...	Sold	
1	96792	160000	0.0	1.0	...	Sold	
2	96771	160000	2.0	1.0	...	Sold	
3	96781	159000	NaN	NaN	...	Sold	
4	96816	160000	2.0	2.0	...	Sold	
5	96771	160000	NaN	NaN	...	Sold	
6	96792	160000	0.0	1.0	...	Sold	
7	96771	160000	2.0	1.0	...	Sold	
8	96816	160000	2.0	2.0	...	Sold	
9	96771	160000	NaN	NaN	...	Sold	

NEXT OPEN HOUSE START TIME	NEXT OPEN HOUSE END TIME	\
----------------------------	--------------------------	---

0	NaN	NaN
1	NaN	NaN
2	NaN	NaN
3	NaN	NaN
4	NaN	NaN
5	NaN	NaN
6	NaN	NaN
7	NaN	NaN
8	NaN	NaN
9	NaN	NaN

URL (SEE <https://www.redfin.com/buy-a-home/comparative-market-analysis> FOR INFO ON PRICING) \

0	<a href="https://www.redfin.com/HI/Waianae/84-770-Kili-...">https://www.redfin.com/HI/Waianae/84-770-Kili-...</a>
1	<a href="https://www.redfin.com/HI/Waianae/85-175-Farri...">https://www.redfin.com/HI/Waianae/85-175-Farri...</a>
2	<a href="https://www.redfin.com/HI/Mountain-View/16-743...">https://www.redfin.com/HI/Mountain-View/16-743...</a>
3	<a href="https://www.redfin.com/HI/Papaikou/Lot-96781/u...">https://www.redfin.com/HI/Papaikou/Lot-96781/u...</a>
4	<a href="https://www.redfin.com/HI/Honolulu/4999-Kahala...">https://www.redfin.com/HI/Honolulu/4999-Kahala...</a>
5	<a href="https://www.redfin.com/HI/Mountain-View/Unknow...">https://www.redfin.com/HI/Mountain-View/Unknow...</a>
6	<a href="https://www.redfin.com/HI/Waianae/85-175-Farri...">https://www.redfin.com/HI/Waianae/85-175-Farri...</a>
7	<a href="https://www.redfin.com/HI/Mountain-View/16-743...">https://www.redfin.com/HI/Mountain-View/16-743...</a>
8	<a href="https://www.redfin.com/HI/Honolulu/4999-Kahala...">https://www.redfin.com/HI/Honolulu/4999-Kahala...</a>
9	<a href="https://www.redfin.com/HI/Mountain-View/Unknow...">https://www.redfin.com/HI/Mountain-View/Unknow...</a>

	SOURCE	MLS#	FAVORITE	INTERESTED	LATITUDE
\					
0	HiCentral MLS	202203123.0	N	Y	21.482810
1	HiCentral MLS	202208655.0	N	Y	21.457467
2	HI Information Service	664218.0	N	Y	19.500440
3	HI Information Service	661374.0	N	Y	19.801533
4	HiCentral MLS	202209911.0	N	Y	21.271372
5	HI Information Service	664972.0	N	Y	19.556521

6	HiCentral MLS	202208655.0	N	Y	21.457467
7	HI Information Service	664218.0	N	Y	19.500440
8	HiCentral MLS	202209911.0	N	Y	21.271372
9	HI Information Service	664972.0	N	Y	19.556521

```

    LONGITUDE
0 -158.203623
1 -158.202598
2 -155.022297
3 -155.108610
4 -157.775244
5 -155.107028
6 -158.202598
7 -155.022297
8 -157.775244
9 -155.107028

```

```
[10 rows x 27 columns]
```

```
df.shape
```

```
(809, 27)
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 809 entries, 0 to 808
```

```
Data columns (total 27 columns):
```

```

#      Column
Non-Null Count  Dtype
---  -
0      SALE TYPE
809 non-null    object
1      SOLD DATE
801 non-null    object
2      PROPERTY TYPE
809 non-null    object
3      ADDRESS
808 non-null    object
4      CITY
809 non-null    object
5      STATE OR PROVINCE
809 non-null    object
6      ZIP OR POSTAL CODE
809 non-null    object

```

```

7 PRICE
809 non-null int64
8 BEDS
552 non-null float64
9 BATHS
533 non-null float64
10 LOCATION
801 non-null object
11 SQUARE FEET
536 non-null float64
12 LOT SIZE
720 non-null float64
13 YEAR BUILT
537 non-null float64
14 DAYS ON MARKET
0 non-null float64
15 $/SQUARE FEET
536 non-null float64
16 HOA/MONTH
355 non-null float64
17 STATUS
801 non-null object
18 NEXT OPEN HOUSE START TIME
0 non-null float64
19 NEXT OPEN HOUSE END TIME
0 non-null float64
20 URL (SEE https://www.redfin.com/buy-a-home/comparative-market-
analysis FOR INFO ON PRICING) 809 non-null object
21 SOURCE
801 non-null object
22 MLS#
801 non-null float64
23 FAVORITE
809 non-null object
24 INTERESTED
809 non-null object
25 LATITUDE
809 non-null float64
26 LONGITUDE
809 non-null float64
dtypes: float64(13), int64(1), object(13)
memory usage: 170.8+ KB

```

```
df.describe()
```

	PRICE	BEDS	BATHS	SQUARE FEET	LOT SIZE
\count	8.090000e+02	552.000000	533.000000	536.000000	7.200000e+02
mean	4.541100e+05	2.235507	1.733583	1128.283582	1.310433e+05

std	6.107641e+05	1.794013	0.929031	783.162290	3.787302e+05
min	5.264000e+03	0.000000	1.000000	29.000000	9.920000e+02
25%	9.000000e+04	1.000000	1.000000	545.500000	8.344750e+03
50%	1.820000e+05	2.000000	1.500000	922.000000	2.150000e+04
75%	7.250000e+05	3.000000	2.000000	1614.000000	8.713100e+04
max	6.750000e+06	18.000000	10.000000	6838.000000	4.295321e+06

	YEAR BUILT	DAYS ON MARKET	\$/SQUARE FEET	HOA/MONTH \
count	537.000000	0.0	536.000000	355.000000
mean	1982.923650	NaN	516.440299	531.873239
std	20.228816	NaN	345.825476	371.519955
min	1920.000000	NaN	17.000000	4.000000
25%	1970.000000	NaN	282.750000	243.000000
50%	1978.000000	NaN	471.500000	498.000000
75%	1996.000000	NaN	656.000000	714.000000
max	2023.000000	NaN	3500.000000	2095.000000

	NEXT OPEN HOUSE START TIME	NEXT OPEN HOUSE END TIME
MLS# \		
count	0.0	0.0
8.010000e+02		
mean	NaN	NaN
7.510187e+07		
std	NaN	NaN
9.737080e+07		
min	NaN	NaN
3.913990e+05		
25%	NaN	NaN
6.610580e+05		
50%	NaN	NaN
6.643950e+05		
75%	NaN	NaN
2.022097e+08		
max	NaN	NaN
2.022237e+08		

	LATITUDE	LONGITUDE
count	809.000000	809.000000
mean	20.459630	-156.524951
std	0.928859	1.368855
min	19.032161	-159.714066
25%	19.535107	-157.841293

```

50%      20.846039 -156.452392
75%      21.309331 -155.103069
max       22.226329 -154.892380

```

```
df.describe(include = ['O'])
```

```

count      SALE TYPE      SOLD DATE      PROPERTY TYPE \
unique              1              76              7
top      PAST SALE  August-25-2022  Single Family Residential
freq              809              25              284

```

```

count      ADDRESS      CITY STATE OR PROVINCE ZIP OR POSTAL
unique              689              51              1
top      4th Ave (awapuhi) Honolulu      HI
freq              5              122              809

```

```

count      LOCATION STATUS \
unique              168      1
top      Waikiki      Sold
freq              62      801

```

```

count      URL (SEE https://www.redfin.com/buy-a-home/comparative-market-
analysis FOR INFO ON PRICING) \
unique              698
top      https://www.redfin.com/HI/Waianae/84-770-Kili-...
freq              2

```

```

count      SOURCE FAVORITE INTERESTED
unique              3          1          1
top      HI Information Service      N          Y
freq              388          809          809

```

```
df['SALE TYPE'].value_counts()
```

```

PAST SALE      809
Name: SALE TYPE, dtype: int64

```

```

df['STATE OR PROVINCE'].value_counts()

HI      809
Name: STATE OR PROVINCE, dtype: int64

df['FAVORITE'].value_counts()

N      809
Name: FAVORITE, dtype: int64

df['INTERESTED'].value_counts()

Y      809
Name: INTERESTED, dtype: int64

df['STATUS'].value_counts()

Sold     801
Name: STATUS, dtype: int64

#drop the columns, which has no data
#drop columns, which has same data

df = df.drop(['SALE TYPE', 'STATE OR PROVINCE', 'FAVORITE',
              'INTERESTED', 'STATUS', 'DAYS ON MARKET', 'NEXT OPEN HOUSE START
              TIME', 'NEXT OPEN HOUSE END TIME'], axis = 1)

df.shape

(809, 19)

df['PROPERTY TYPE'].value_counts().reset_index()

   index  PROPERTY TYPE
0  Single Family Residential      284
1           Vacant Land      261
2           Condo/Co-op      256
3  Multi-Family (5+ Unit)         3
4  Multi-Family (2-4 Unit)         2
5           Townhouse         2
6              Other         1

sns.set_style("whitegrid")
plt.figure(figsize = (8,5))
plt.title('Distribution of PROPERTY TYPE', fontsize=18,
fontweight='bold')
eda_percentage = df['PROPERTY TYPE'].value_counts(normalize =
True).rename_axis('PROPERTY TYPE').reset_index(name = 'Percentage')

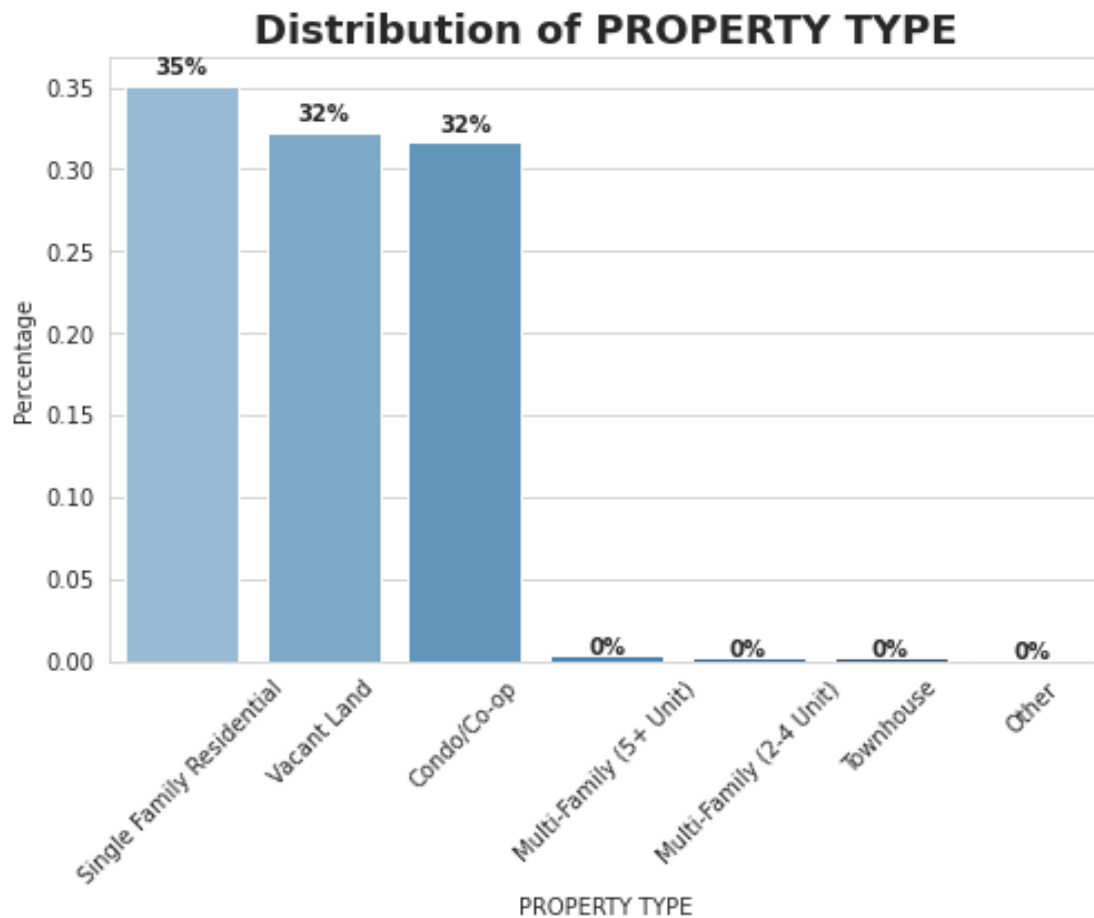
ax = sns.barplot(x = 'PROPERTY TYPE', y = 'Percentage', data =
eda_percentage.head(10), palette='Blues_d')
for p in ax.patches:
    width = p.get_width()

```

```

height = p.get_height()
x, y = p.get_xy()
ax.annotate(f'{height:.0%}', (x + width/2, y + height*1.02),
ha='center', fontweight='bold')
plt.setp(ax.get_xticklabels(), rotation=45);

```



```

fig, ax = plt.subplots(figsize=(15, 8))
chart = sns.countplot(x="CITY", data=df, ax=ax, palette='Blues_d',
                      order = df['CITY'].value_counts().index)
chart.set_xticklabels(chart.get_xticklabels(), rotation=90)
plt.show()

```





```

663759.0      2
202210561.0    1
202213665.0    1
202211939.0    1
664274.0       1
202211701.0    1
Name: MLS#, Length: 690, dtype: int64

```

```


sns.set_style("whitegrid")
plt.figure(figsize = (8,5))
plt.title('Distribution of BEDS', fontsize=18, fontweight='bold')
eda_percentage = df['BEDS'].value_counts(normalize =
True).rename_axis('BEDS').reset_index(name = 'Percentage')

```

```

ax = sns.barplot(x = 'BEDS', y = 'Percentage', data = eda_percentage,
palette='Blues_d', order =
eda_percentage['BEDS'].value_counts().index)
for p in ax.patches:
    width = p.get_width()
    height = p.get_height()
    x, y = p.get_xy()
    ax.annotate(f'{height:.0%}', (x + width/2, y + height*1.02),
ha='center', fontweight='bold')
plt.setp(ax.get_xticklabels(), rotation=45);

```



eda\_percentage

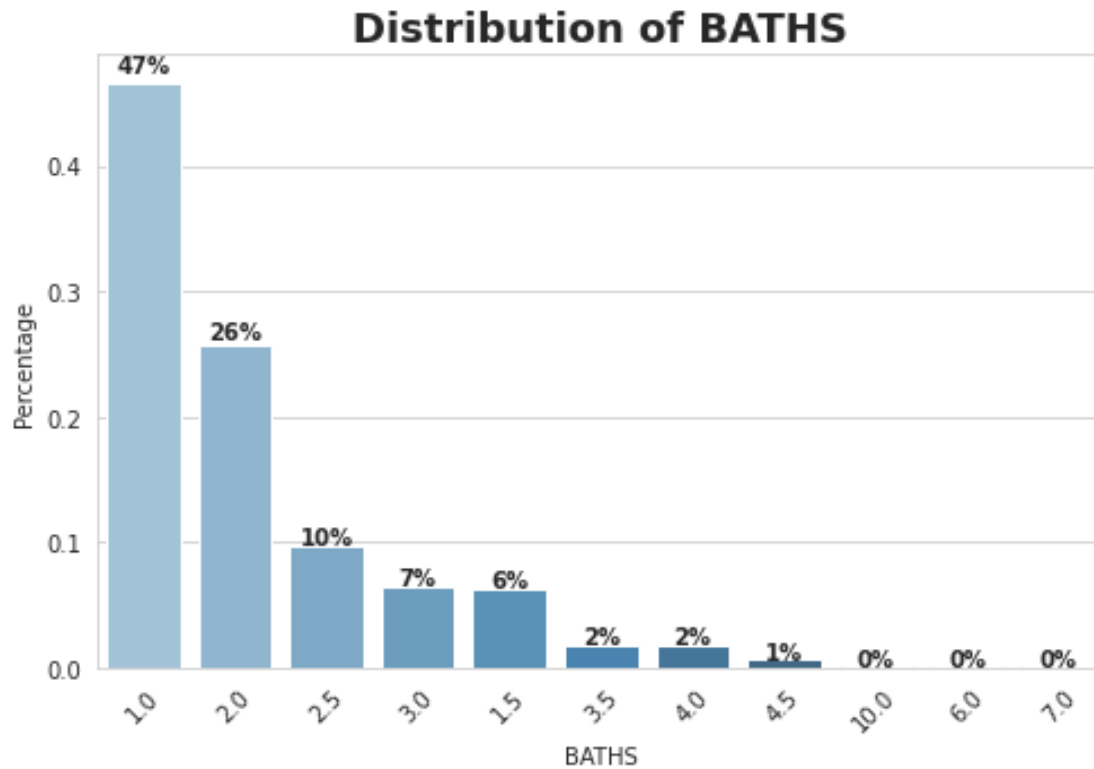
	BEDS	Percentage
0	3.0	0.275362
1	1.0	0.193841
2	2.0	0.184783
3	0.0	0.170290
4	4.0	0.125000
5	5.0	0.027174
6	6.0	0.007246
7	11.0	0.005435
8	9.0	0.003623
9	8.0	0.003623
10	7.0	0.001812
11	18.0	0.001812

```

sns.set_style("whitegrid")
plt.figure(figsize = (8,5))
plt.title('Distribution of BATHS', fontsize=18, fontweight='bold')
eda_percentage = df['BATHS'].value_counts(normalize =
True).rename_axis('BATHS').reset_index(name = 'Percentage')

ax = sns.barplot(x = 'BATHS', y = 'Percentage', data = eda_percentage,
palette='Blues_d', order =
eda_percentage['BATHS'].value_counts().index)
for p in ax.patches:
    width = p.get_width()
    height = p.get_height()
    x, y = p.get_xy()
    ax.annotate(f'{height:.0%}', (x + width/2, y + height*1.02),
ha='center', fontweight='bold')
plt.setp(ax.get_xticklabels(), rotation=45);

```



```
df['BATHS'].value_counts().reset_index()
```

	index	BATHS
0	1.0	248
1	2.0	137
2	2.5	52
3	3.0	35
4	1.5	34
5	3.5	10
6	4.0	10
7	4.5	4
8	10.0	1
9	6.0	1
10	7.0	1

```
# forward-fill
```

```
df = df.fillna(method='ffill')
```

```
df.isnull().sum()
```

```
SOLD DATE
```

```
0
```

```
PROPERTY TYPE
```

```
0
```

```
ADDRESS
```

```
0
```

```
CITY
```

```

0
ZIP OR POSTAL CODE
0
PRICE
0
BEDS
0
BATHS
0
LOCATION
0
SQUARE FEET
0
LOT SIZE
0
YEAR BUILT
0
$/SQUARE FEET
0
HOA/MONTH
0
URL (SEE https://www.redfin.com/buy-a-home/comparative-market-analysis
FOR INFO ON PRICING) 0
SOURCE
0
MLS#
0
LATITUDE
0
LONGITUDE
0
dtype: int64

df['YEAR BUILT'] = df['YEAR BUILT'].astype("int")

labels = ['1920', '1930', '1940', '1950', '1960', '1970', '1980',
'1990', '2000', '2010', '2020']
bins = [ 1920, 1930, 1940, 1950, 1960, 1970, 1980, 1990, 2000, 2010,
2020, 2030 ]
df['year_bin'] = pd.cut(df['YEAR BUILT'] , bins=bins, labels=labels,
include_lowest=True)

df['year_bin'].nunique()

11


fig, ax = plt.subplots(figsize=(15, 8))
chart = sns.countplot(x="year_bin", data=df, ax=ax, palette='Blues_d',

```

```

order = df['year_bin'].value_counts().index
chart.set_xticklabels(chart.get_xticklabels(), rotation=45)
plt.show()

```



```
df.columns
```

```

Index(['SALE TYPE', 'SOLD DATE', 'PROPERTY TYPE', 'ADDRESS', 'CITY',
      'STATE OR PROVINCE', 'ZIP OR POSTAL CODE', 'PRICE', 'BEDS',
      'BATHS',
      'LOCATION', 'SQUARE FEET', 'LOT SIZE', 'YEAR BUILT', '$/SQUARE
FEET',
      'HOA/MONTH', 'STATUS',
      'URL (SEE https://www.redfin.com/buy-a-home/comparative-market-
analysis FOR INFO ON PRICING)',
      'SOURCE', 'MLS#', 'FAVORITE', 'INTERESTED', 'LATITUDE',
      'LONGITUDE'],
      dtype='object')

```


```

cts_variables = ['PRICE', 'SQUARE FEET', 'LOT SIZE', '$/SQUARE FEET',
                 'HOA/MONTH', 'LATITUDE', 'LONGITUDE']

```

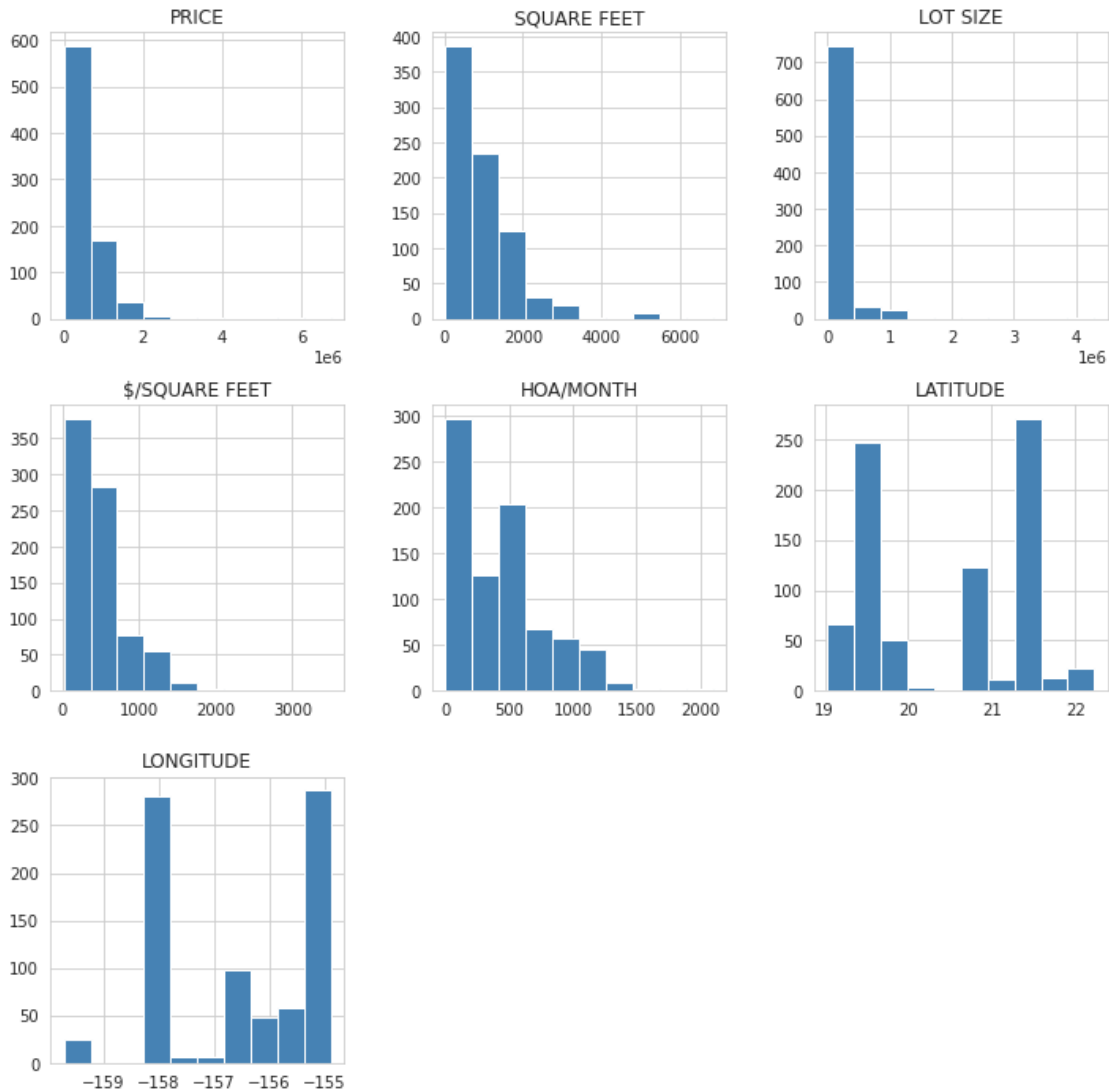
```
sns.pairplot(df[cts_variables], corner=True)
```

```
<seaborn.axisgrid.PairGrid at 0x7f86ae801050>
```



*# Distribution of continuous Features of the Dataset*

```
distribution = df[cts_variables].hist( linewidth = 1.0, color =
'steelblue')
fig = plt.gcf()
fig.set_size_inches(12,12)
plt.show()
```



```
pd.set_option('display.float_format', lambda x: '%.5f' % x)
```

```
df['PRICE'].value_counts()
```

```
165000    14
175000    12
150000    12
140000    12
200000    11
..
301000     1
455000     1
669000     1
635000     1
829000     1
Name: PRICE, Length: 412, dtype: int64
```




```
plt.rcParams['figure.figsize'] = [8, 5]
sns.set(style = 'white')
plt.title('Distribution of PRICE', fontsize=18, fontweight='bold')
sns.distplot(df['PRICE'], color = "steelblue")
plt.ylabel("Distribution", fontsize = 15)
plt.xlabel("PRICE", fontsize = 15)
plt.margins(x = 0)

print ("The maximum PRICE is", df['PRICE'].max())
print ("The minimum PRICE is", df['PRICE'].min())

/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619:
FutureWarning: `distplot` is a deprecated function and will be removed
in a future version. Please adapt your code to use either `displot` (a
figure-level function with similar flexibility) or `histplot` (an
axes-level function for histograms).
  warnings.warn(msg, FutureWarning)

The maximum PRICE is 6750000
The minimum PRICE is 5264
```



```
df['PRICE'] = df['PRICE'].astype("int")

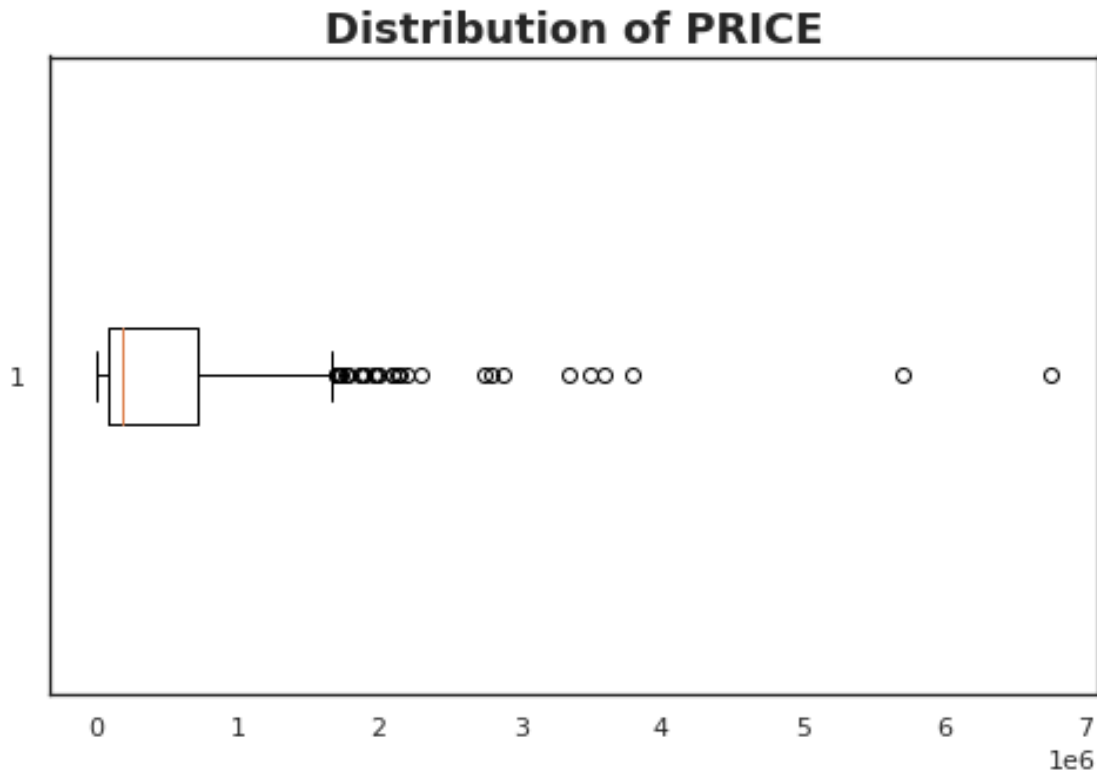
from matplotlib import pyplot as plt, ticker as mticker

fig, ax = plt.subplots(1, 1)
# Creating plot
plt.boxplot(df['PRICE'],
```

```

        vert = 0,
    )
plt.title('Distribution of PRICE', fontsize=18, fontweight='bold')
plt.show()


```



```

fig, ax = plt.subplots(1, 1)
# Creating plot
plt.boxplot(df['PRICE'],
            vert = 0,
            )
plt.title('Distribution of PRICE', fontsize=18, fontweight='bold')
ax = plt.gca()
ax.set_xscale('log')
ax.set_yscale('log')
ax.xaxis.set_minor_formatter(mticker.ScalarFormatter())
plt.show()

```



```
df['PRICE'].describe().reset_index()
```

index	PRICE
0 count	8.090000e+02
1 mean	4.541100e+05
2 std	6.107641e+05
3 min	5.264000e+03
4 25%	9.000000e+04
5 50%	1.820000e+05
6 75%	7.250000e+05
7 max	6.750000e+06

```
labels = ['<100K', '100K-500K', '500K-1M', '1M+']
bins = [ 5264, 100000, 500000, 1000000, 6750000]
df['PRICE_bin'] = pd.cut(df['PRICE'] , bins=bins, labels=labels,
include_lowest=True)
```

```
df['PRICE_bin'].nunique()
```

```
4
```

```
df['PRICE_bin'].value_counts().reset_index()
```


index	PRICE_bin
0 100K-500K	324
1 <100K	225
2 500K-1M	155
3 1M+	105

```

sns.set_style("whitegrid")
plt.figure(figsize = (8,5))
plt.title('Distribution of PRICE_bin', fontsize=18, fontweight='bold')
eda_percentage = df['PRICE_bin'].value_counts(normalize =
True).rename_axis('PRICE_bin').reset_index(name = 'Percentage')

ax = sns.barplot(x = 'PRICE_bin', y = 'Percentage', data =
eda_percentage, palette='Blues_d', order =
eda_percentage['PRICE_bin'].value_counts().index)
for p in ax.patches:
    width = p.get_width()
    height = p.get_height()
    x, y = p.get_xy()
    ax.annotate(f'{height:.0%}', (x + width/2, y + height*1.02),
ha='center', fontweight='bold')
plt.setp(ax.get_xticklabels(), rotation=45);

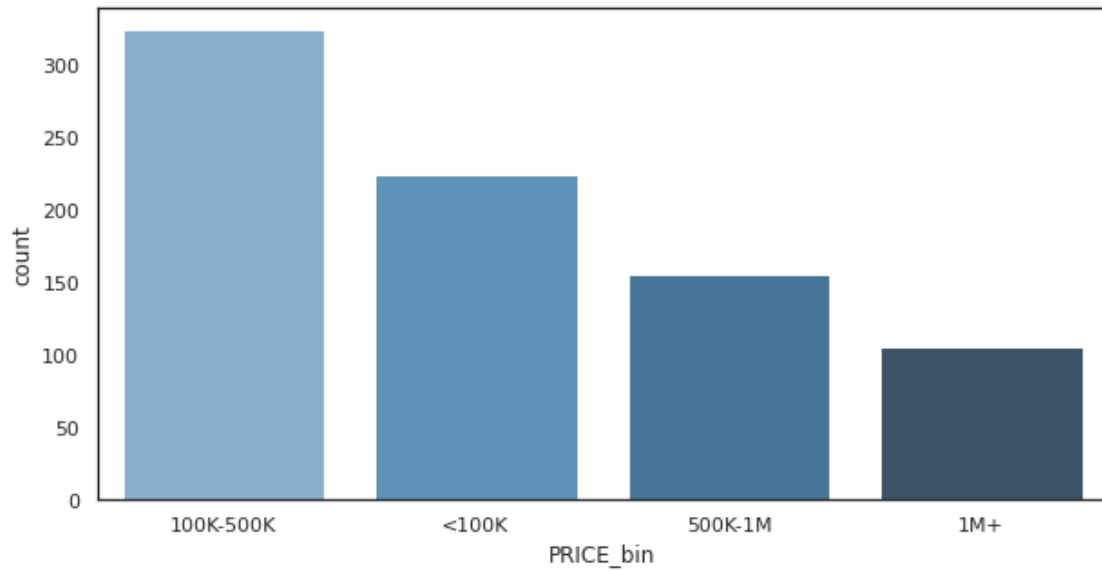
```



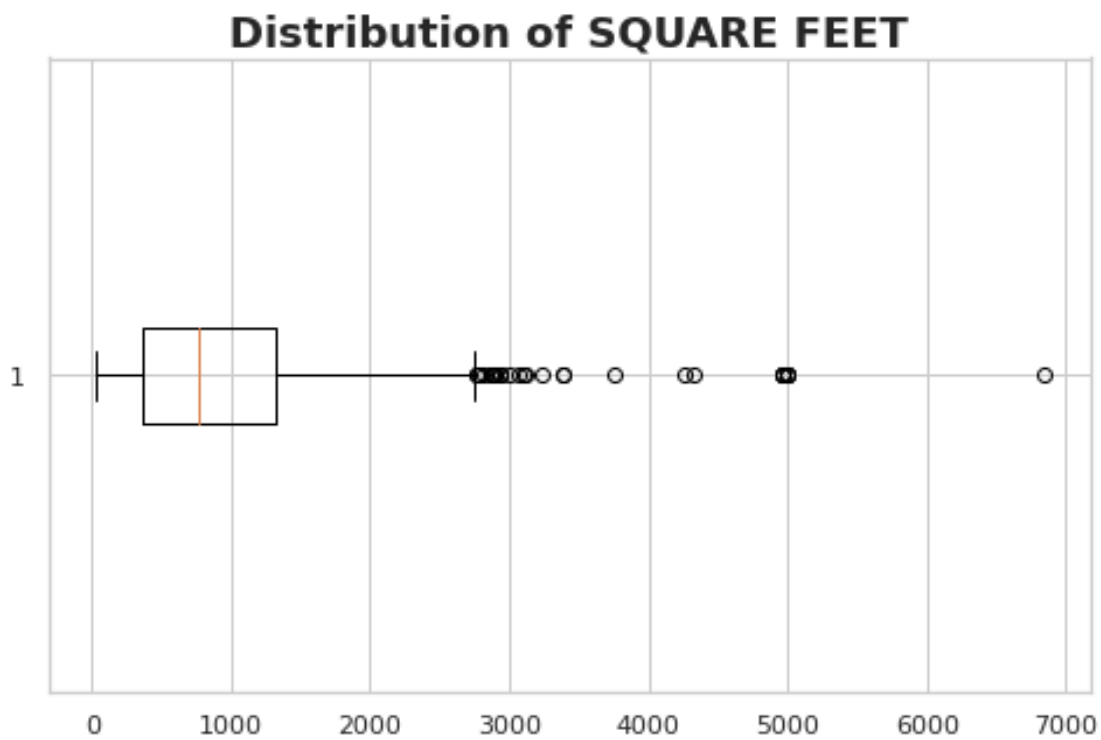
```

fig, ax = plt.subplots(figsize=(10, 5))
chart = sns.countplot(x="PRICE_bin", data=df, ax=ax,
palette='Blues_d',
order = df['PRICE_bin'].value_counts().index)
chart.set_xticklabels(chart.get_xticklabels(), rotation=0)
plt.show()

```




```
fig, ax = plt.subplots(1, 1)
# Creating plot
plt.boxplot(df['SQUARE FEET'],
            vert = 0,
            )
plt.title('Distribution of SQUARE FEET', fontsize=18,
          fontweight='bold')
plt.show()
```



```
cts_variables2 = ['SQUARE FEET', '$/SQUARE FEET', 'HOA/MONTH',
                  'LATITUDE', 'LONGITUDE']
```

```
distribution = df[cts_variables2].boxplot(color = 'steelblue')
fig = plt.gcf()
fig.set_size_inches(12,12)
plt.show()
```



```
df.columns
```

```
Index(['SALE TYPE', 'SOLD DATE', 'PROPERTY TYPE', 'ADDRESS', 'CITY',
      'STATE OR PROVINCE', 'ZIP OR POSTAL CODE', 'PRICE', 'BEDS',
      'BATHS', 'LOCATION', 'SQUARE FEET', 'LOT SIZE', 'YEAR BUILT', 'DAYS ON
MARKET', '$/SQUARE FEET', 'HOA/MONTH', 'STATUS', 'NEXT OPEN HOUSE START
TIME', 'NEXT OPEN HOUSE END TIME',
      'URL (SEE https://www.redfin.com/buy-a-home/comparative-market-
```


```
analysis FOR INFO ON PRICING)',
      'SOURCE', 'MLS#', 'FAVORITE', 'INTERESTED', 'LATITUDE',
      'LONGITUDE',
      'bin_cut_manual', 'PRICE_bin'],
      dtype='object')
```

What is the most expensive location in Hawaii?

```
df_temp =
df.groupby(['CITY']).PRICE.mean().sort_values(ascending=False)[:10]
df_temp.head()
```

```
CITY
Kilauea    4425000.00000
Hanalei    3175000.00000
Paia       2804000.00000
Kula       2800000.00000
Kailua     1980000.00000
Name: PRICE, dtype: float64
```

```
df.plot(kind="scatter", x='LONGITUDE', y='LATITUDE', alpha=0.4,
c=df['PRICE'], s=10,
      cmap=plt.get_cmap('jet'), figsize=(12,8));
```




```
sns.lmplot(
    "LONGITUDE", "LATITUDE", data=df, hue="PRICE_bin", fit_reg=False,
    size=7
);
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/regression.py:581:  
UserWarning: The `size` parameter has been renamed to `height`; please  
update your code.
```


```
warnings.warn(msg, UserWarning)
```

<Figure size 1728x288 with 0 Axes>



```
fig, ax = plt.subplots(figsize=(25, 10))  
sns.boxplot(x='CITY', y='PRICE', data=df, ax=ax)  
plt.setp(ax.get_xticklabels(), rotation=90);
```





1e6 = 1\*10<sup>6</sup>

```
df_new =
df.groupby(['CITY']).PRICE.mean().sort_values(ascending=False).reset_index()
df_new.head(10)
```

	CITY	PRICE
0	Kilauea	4425000.00000
1	Hanalei	3175000.00000
2	Paia	2804000.00000
3	Kula	2800000.00000
4	Kailua	1980000.00000
5	Princeville	1463583.33333
6	Haiku	1427090.81818
7	Kahuku	1379057.00000
8	Makawao	1338416.66667
9	Kaaawa	1277500.00000

df.columns


```
Index(['SALE TYPE', 'SOLD DATE', 'PROPERTY TYPE', 'ADDRESS', 'CITY',
      'STATE OR PROVINCE', 'ZIP OR POSTAL CODE', 'PRICE', 'BEDS',
      'BATHS',
      'LOCATION', 'SQUARE FEET', 'LOT SIZE', 'YEAR BUILT', 'DAYS ON
MARKET',
      '$/SQUARE FEET', 'HOA/MONTH', 'STATUS', 'NEXT OPEN HOUSE START
TIME',
      'NEXT OPEN HOUSE END TIME',
      'URL (SEE https://www.redfin.com/buy-a-home/comparative-market-
analysis FOR INFO ON PRICING)',
      'SOURCE', 'MLS#', 'FAVORITE', 'INTERESTED', 'LATITUDE',
      'LONGITUDE',
      'bin_cut_manual', 'PRICE_bin'],
      dtype='object')
```

```
sns.lmplot(
    'SQUARE FEET', 'LOT SIZE', data=df, hue="PRICE_bin",
    fit_reg=False, size=7
);
```

/usr/local/lib/python3.7/dist-packages/seaborn/\_decorators.py:43:  
FutureWarning: Pass the following variables as keyword args: x, y.  
From version 0.12, the only valid positional argument will be `data`,  
and passing other arguments without an explicit keyword will result in  
an error or misinterpretation.


FutureWarning  
/usr/local/lib/python3.7/dist-packages/seaborn/regression.py:581:  
UserWarning: The `size` parameter has been renamed to `height`; please  
update your code.

warnings.warn(msg, UserWarning)




LOT size indicates the size of the piece of land where the property is situated.


```
fig, ax = plt.subplots(figsize=(15, 8))
sns.countplot(x="PROPERTY TYPE", hue="PRICE_bin", data=df)
plt.setp(ax.get_xticklabels(), rotation=90);
```



```
fig, ax = plt.subplots(figsize=(15, 8))
sns.countplot(x="BEDS", hue="PRICE_bin", data=df)
plt.setp(ax.get_xticklabels(), rotation=90);
```




```
fig, ax = plt.subplots(figsize=(15, 8))
sns.countplot(x="BATHS", hue="PRICE_bin", data=df)
plt.setp(ax.get_xticklabels(), rotation=0);
```




```
sns.lmplot(
    'BEDS', 'BATHS', data=df, hue="PRICE_bin", fit_reg=False, height=7
);
```

/usr/local/lib/python3.7/dist-packages/seaborn/\_decorators.py:43:  
FutureWarning: Pass the following variables as keyword args: x, y.  
From version 0.12, the only valid positional argument will be `data`,  
and passing other arguments without an explicit keyword will result in  
an error or misinterpretation.  
FutureWarning



```
fig, ax = plt.subplots(figsize=(15, 8))
sns.countplot(x="year_bin", hue="PRICE_bin", data=df)
plt.setp(ax.get_xticklabels(), rotation=90);
```



```
df.groupby(['year_bin']).PRICE.mean().sort_values(ascending=False).reset_index()
```

	year_bin	PRICE
0	1930	1.105000e+06
1	1990	9.357390e+05
2	2000	8.754926e+05
3	2020	8.399239e+05
4	1920	6.587110e+05
5	1980	4.640043e+05
6	1940	4.267059e+05
7	2010	4.048218e+05
8	1970	3.479158e+05
9	1950	3.428571e+05
10	1960	2.475584e+05

```
df.groupby(['YEAR BUILT']).PRICE.mean().sort_values(ascending=False).reset_index().head(10)
```

	YEAR BUILT	PRICE
0	1996	2759333.33333
1	1955	2200000.00000
2	2004	1890000.00000
3	2019	1504750.00000
4	1965	1371666.66667
5	1929	1260000.00000
6	1999	1247541.66667
7	1949	1247500.00000
8	2003	1244850.00000
9	2018	1239000.00000

```
df['SOLD Month'] = df['SOLD DATE'].str.split('-').str[0]
df['SOLD Month']
```

0	September
1	August
2	September
3	September
4	August
	...
804	September
805	October
806	October
807	September
808	September

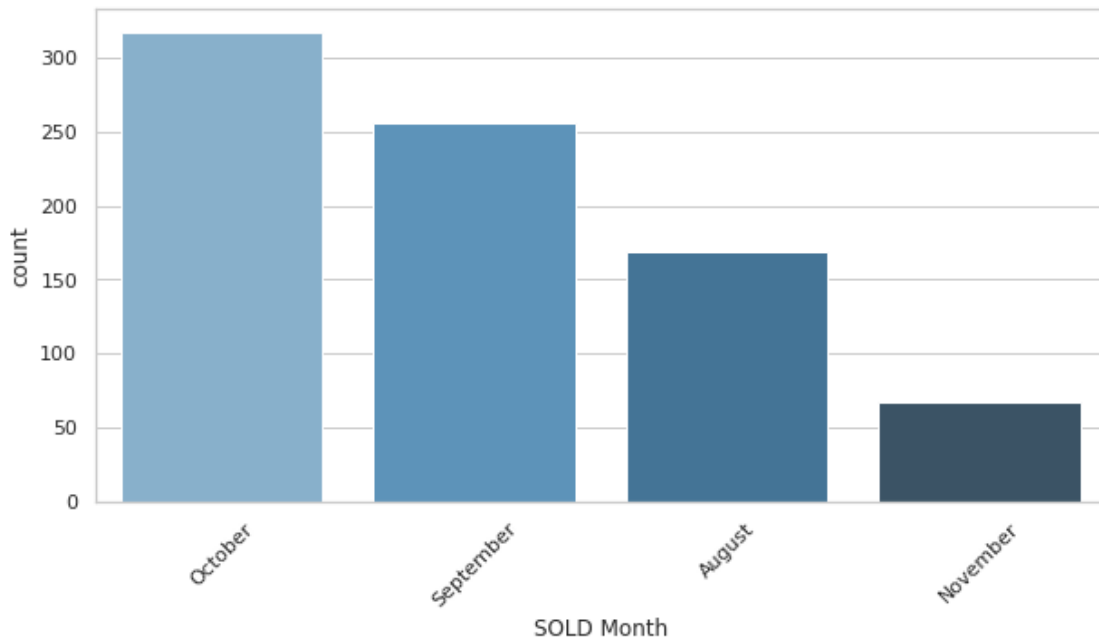
Name: SOLD Month, Length: 809, dtype: object

```
fig, ax = plt.subplots(figsize=(10, 5))
chart = sns.countplot(x='SOLD Month', data=df, ax=ax,
palette='Blues_d',
```

```

        order = df['SOLD Month'].value_counts().index)
chart.set_xticklabels(chart.get_xticklabels(), rotation=45)
plt.show()

```



```

df.groupby(['SOLD
Month']).PRICE.mean().sort_values(ascending=False).reset_index()


```

	SOLD Month	PRICE
0	November	608129.85075
1	October	476900.89590
2	September	454972.25000
3	August	348992.92308


```

fig, ax = plt.subplots(figsize=(15, 10))
sns.boxplot(x='SOLD Month', y='PRICE', data=df, ax=ax)
plt.setp(ax.get_xticklabels(), rotation=90);

```



```
fig, ax = plt.subplots(figsize=(15, 8))
sns.countplot(x="SOLD Month", hue="PRICE_bin", data=df)
plt.setp(ax.get_xticklabels(), rotation=90);
```




```
df.columns
```



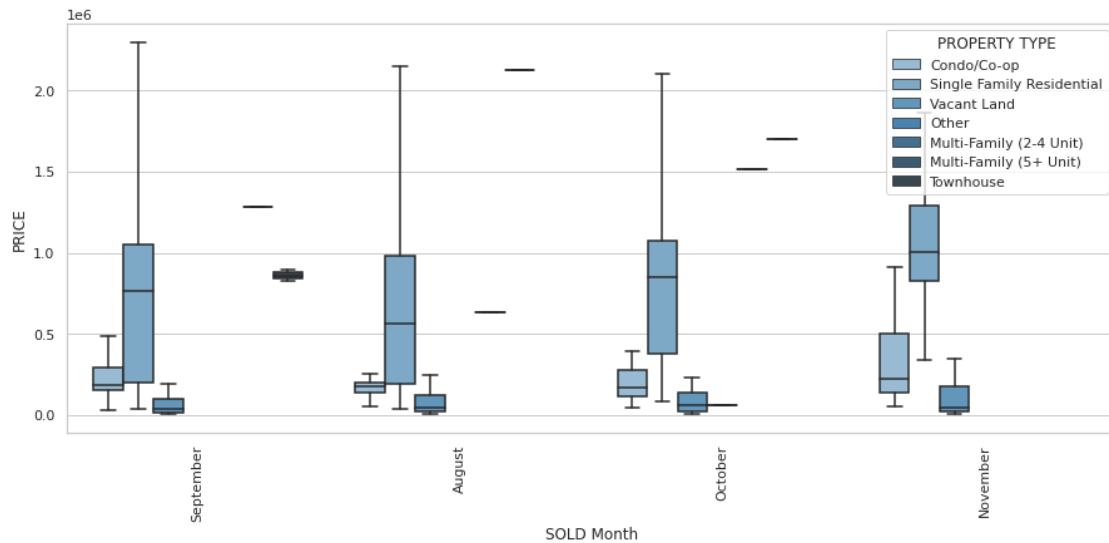
```
Index(['SOLD DATE', 'PROPERTY TYPE', 'ADDRESS', 'CITY', 'ZIP OR POSTAL
CODE',
      'PRICE', 'BEDS', 'BATHS', 'LOCATION', 'SQUARE FEET', 'LOT
SIZE',
      'YEAR BUILT', '$/SQUARE FEET', 'HOA/MONTH',
      'URL (SEE https://www.redfin.com/buy-a-home/comparative-market-
analysis FOR INFO ON PRICING)',
      'SOURCE', 'MLS#', 'LATITUDE', 'LONGITUDE', 'bin_cut_manual',
      'PRICE_bin', 'year_bin', 'SOLD Month'],
      dtype='object')
```

```
fig, ax = plt.subplots(figsize=(15, 8))
sns.countplot(x="SOLD Month", hue="PROPERTY TYPE", data=df)
plt.setp(ax.get_xticklabels(), rotation=90);
```



```
def boxplot_variation(feature1, feature2, feature3, width=16):
    fig, ax1 = plt.subplots(ncols=1, figsize=(width,6))
    s = sns.boxplot(ax = ax1, x=feature1, y=feature2, hue=feature3,
                    data=df, palette="Blues_d", showfliers=False)
    s.set_xticklabels(s.get_xticklabels(), rotation=90)
    plt.show();
```

```
boxplot_variation('SOLD Month', 'PRICE', 'PROPERTY TYPE', 15)
```



```
plt.figure(figsize = (12,10))
plt.title("Correlation between different features of the dataset",
          fontsize = 18, fontweight = 'bold')
sns.heatmap(df[['PRICE', 'SQUARE FEET', 'LOT SIZE', '$/SQUARE FEET',
                'HOA/MONTH',
                'LATITUDE', 'LONGITUDE', 'BEDS', 'BATHS']].corr(), cmap =
'Blues', annot = True)
plt.xticks(fontsize=12, rotation = 90)
plt.yticks(fontsize=12, rotation = 0)
plt.legend(fontsize=12)
plt.show()
```

WARNING:matplotlib.legend:No handles with labels found to put in legend.

Correlation between different features of the dataset

